

# Deep Learning for CIFAR-10 Image Classification

Exploring Convolutional Neural Networks  
(CNNs) and Transfer Learning

September 8, 2024

# Project Overview



CIFAR-10 Image Classification using CNN and Transfer Learning



Objective: Classify images from CIFAR-10 into 10 categories using deep learning.

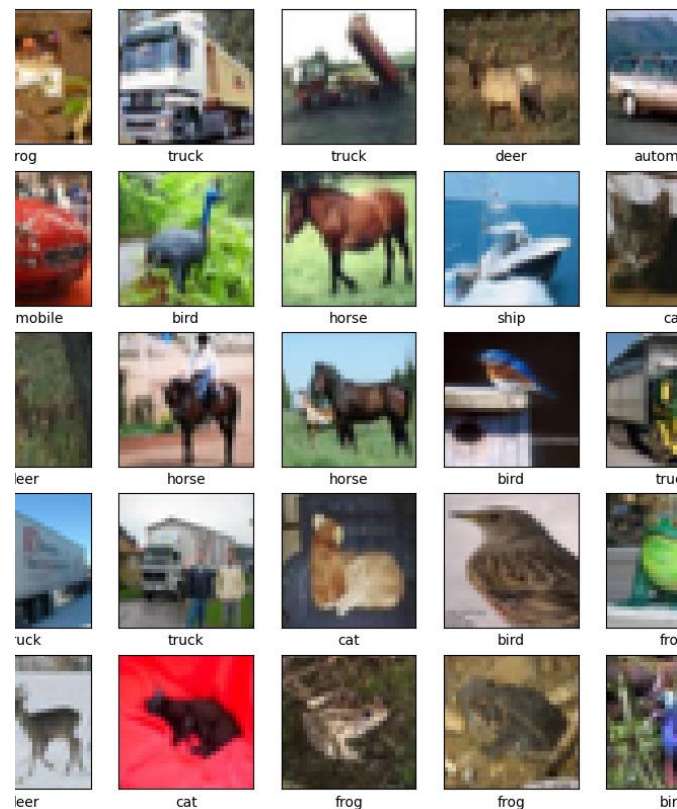


Dataset: 60,000 images (32x32) across 10 categories (airplane, car, bird, etc.).

# Problem Statement

Challenge: Accurately classify low-resolution images with high variability.

Goal: Build a model that generalizes well, achieving high accuracy on unseen images.



# ML Approach



CNN Architecture: Multiple convolutional layers with batch normalization and dropout.



Transfer Learning: Pre-trained models (ResNet18, Xception) to improve classification accuracy.



Techniques: Data normalization, Image augmentation (horizontal flipping), Hyperparameter tuning (Grid Search).

# Exploratory Data Analysis (EDA)



Class Distribution: Balanced across the 10 categories.



Image Characteristics: RGB channels, pixel intensity distributions.



Visualization: PCA and t-SNE for class separation, Edge detection using Canny.

# Exploratory Data Analysis (EDA)

Example  
training  
images



frog



truck



truck



deer



automobile



automobile



bird



horse



ship



cat



deer



horse



horse



bird



truck



truck



truck



cat



bird



frog



deer




cat



frog



frog



bird

## Exploratory Data Analysis (EDA)

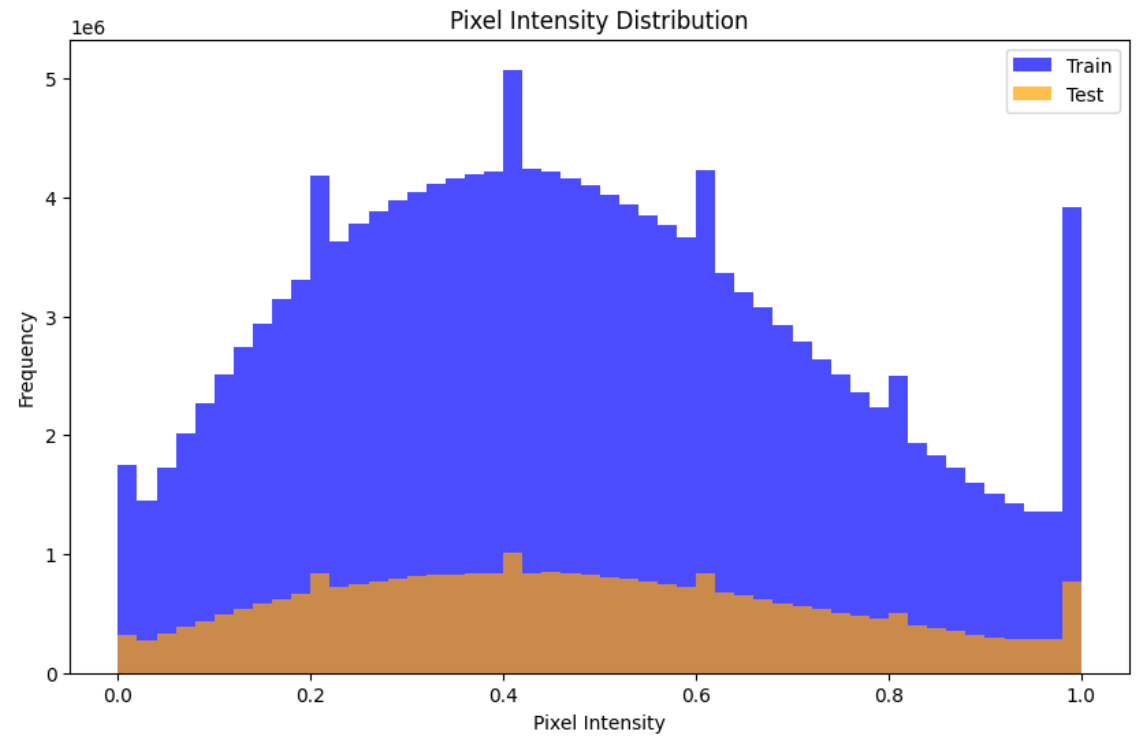
Training data: 10 classes  
Balanced at 5000 images each

	Class	Count
0	airplane	5000
1	automobile	5000
2	bird	5000
3	cat	5000
4	deer	5000
5	dog	5000
6	frog	5000
7	horse	5000
8	ship	5000
9	truck	5000



# Exploratory Data Analysis (EDA)


## Pixel Intensity Distribution






# Exploratory Data Analysis (EDA)

## PCA



# Exploratory Data Analysis (EDA)

t-SNE



# Model Architecture



CNN Model: Convolutional layers, Batch normalization, Dropout to prevent overfitting.



Transfer Learning: Fine-tuned Xception and ResNet18.

# Training Results



CNN Performance: Best overall Test accuracy of 82.5%.



Xception Performance: After tuning, achieved ~80% test accuracy.




ResNet50: Lower performance (38.13%) on CIFAR-10.

# Visualization of Model Performance

Accuracy & Loss Curves: Training and validation accuracy indicate overfitting.

CNN with 5  
convolutional layers,  
50 epochs

# Training Results



# CNN Model Architecture Summary

## **Input Layer:**

- Shape: (32, 32, 3) RGB image

## **Convolutional Blocks (x5):**

- Conv2D: Detects image patterns (32-256 filters)
- BatchNormalization: Stabilizes training
- MaxPooling (2x2): Reduces dimensions
- Dropout (25%): Prevents overfitting

## **GlobalAveragePooling2D:**

- Converts feature maps to a 1D vector

## **Dense Layers:**

- Dense (64 units, ReLU): Learns complex patterns
- Dropout (50%): Regularization
- Dense (10 units): Output layer for classification



## **Compilation:**

- Optimizer: Adam
- Loss: SparseCategoricalCrossentropy
- Metric: Accuracy

## **Training Results:**



- Accuracy: Training vs Validation
- Test Accuracy: Final accuracy





# Training Results

CNN with Random  
Horizontal Flip, 500 epochs



# Training Results

Xception (pretrained) with  
best hyper parameters,  
500 epochs

# Optimization & Hyperparameter Tuning



Hyperparameters: Learning rate, batch size, optimizer choice.



Grid Search: Tested different combinations to maximize accuracy.

# Optimization & Hyperparameter Tuning

Ran each of these 4 pretrained models:

- ResNet50 and ResNet18
- EfficientNetB0
- InceptionV3
- Xception

# Optimization & Hyperparameter Tuning

---

Hyperparameter grid – best parameters found in **bold** below for pretrained Xception

---

lrs = [**0.001**, 0.01, 0.1]

---

momentums = [**0.8**, 0.9]

---

weight\_decays = [**0**, 1e-4]

---

batch\_sizes = [**64**, 128]

---

optimizers = ['SGD', '**Adam**']

---


# Conclusion



Key Takeaways: CNNs are effective for image classification.



Transfer learning with Xception shows promising results.



Future Work: More data augmentation, regularization, deeper models.

# Project Links:

- GitHub Link:  
<https://github.com/LEBLAPI1/ImageClassifier-CNN-ResNet-Xception/>
- Notebook Link (Google Colab):  
[https://colab.research.google.com/drive/1muNErLFUuOlqb1wMnoUJ\\_zdv6eGhyTuW?usp=sharing](https://colab.research.google.com/drive/1muNErLFUuOlqb1wMnoUJ_zdv6eGhyTuW?usp=sharing)
- Data set:  
<https://www.tensorflow.org/datasets/catalog/cifar10>